

Ziyu Xu – Full-text redline against original resume

~~Red strikeout: deleted~~ Blue underline: added

Ziyu_Xu_D3

Original baseline: CMU_Ziyu_original.tex. All visible text is compared in PDF reading order; moved content appears as deletion and insertion. Layout-only changes are described in changes.md.

Ziyu (Rick) Xu Robotics Institute, Carnegie Mellon University Research Interests: ~~Robot Foundation~~ Vision-Language Models, ~~Real-World Policy Adaptation~~, Streaming Video, and ~~Learning-Augmented Control~~ Efficient Inference +1 7348005258 | rickx@andrew.cmu.edu

Education

Carnegie Mellon University (CMU) Aug. 2026 – Expected 2028 Master of Science in Robotics (MSR) Pittsburgh, PA
University of Michigan, Ann Arbor (UM) Oct. 2024 – May 2026 B.S.E. in Mechanical Engineering (Robotics Track);
Minor in Computer Science GPA: 3.85/4.0 (Top 5%) Core Coursework: Robot Learning, RL, Embedded Systems, Optimal Control, Convex Optimization
Shanghai Jiao Tong University (SJTU) Sep. 2022 – Jul. 2026 B.S.E. in Electrical and Computer Engineering GPA: 3.56/4.0 (Top 20%) Core Coursework: Algorithms, Mathematical Foundations, Probability Theory, Signals and Systems

Research Experiences and Internships

~~Unitree Robotics, R&D Department, Embodied AI Team May 2026 – Present Research Intern, Humanoid Robot Foundation Models Hangzhou, China Closed-Loop Real-Robot Infrastructure for Humanoid Foundation-Model Policies: Built and validated an end- to-end policy infrastructure on Unitree H2/G1 humanoids, integrating teleoperation data collection, VLA/WAM and whole-body policy adapters, rollout logging and evaluation, controller interfaces, and policy fine-tuning. Enabled launch with one command and closed-loop rollouts across platforms, allowing policies and controllers to be integrated without modifying platform-specific launch workflows. Embodiment VLA Adaptation & Post-Training: Conducted a structured survey of representative robot VLA, WAM, and physical-AI policies, especially work from NVIDIA GEAR Lab. Deployed and fine-tuned Isaac GR00T/OpenPI VLA policies with SONIC whole-body control on Unitree humanoids using real-robot rollout traces. Currently developing a phase-aware adaptation method that aligns semantic task progress and execution timing across embodiments for policy transfer and rollout-driven post-training. Reproducible Real-World VLA Benchmark for Humanoid Manipulation: Implemented an 8-task, 24-scene benchmark with green-screen background replacement to evaluate foundation models under varied visual con- ditions. Built a PICO-based scene-digitization and AR-guided reconstruction pipeline that captured shelf-object layouts as reusable digital scene templates, and registered target placements back into the physical workspace, achieving 2.5 cm median placement error. Shanghai Artificial Intelligence Lab & IRMV, SJTU Apr. 2026 – Present Research Assistant, Directed by Prof. Junchi Yan and Prof. Hesheng Wang Shanghai, China Context-Aware Re-encoding for Iterative Visual KV Compression: Identified representation Studied errors caused by from reusing retained visual KV entries after their spatial neighbors neighboring tokens are removed, and showed that these errors com- pound across repeated compression removed in multi-turn and streaming inference. Developed a context-aware re-encoding method that recomputes each retained set in its new context at the original positions, positions and benchmarked it against compared with inheritance and restoration methods under matched memory and runtime budgets. Unitree Robotics, R&D Department, Embodied AI Team May 2026 – Present Research Intern, Humanoid Robot Foundation Models Hangzhou, China VLA Adaptation and Post-Training: Deployed and fine-tuned Isaac GR00T/OpenPI policies with SONIC whole- body control using real-robot rollout traces. Developing phase-aware adaptation across embodiments by aligning semantic task progress and execution timing. Reproducible Evaluation: Implemented an 8-task, 24-scene manipulation benchmark with background variation. Built PICO-based scene digitization and AR-guided reconstruction with reusable layouts, achieving 2.5 cm median placement error. Surgical Computer Vision Detection, SJTU Apr. 2024 – Dec. 2024 Research Assistant, Directed by Prof. Yutong Ban Shanghai, China Surgical Video Segmentation: Evaluated Segment Anything Model variants on surgical frames, examining do- main shift and temporal mask stability around instruments, tissues, and interaction regions. Ziyu (Rick) Xu Research experience continued Intelligent Control Lab, Carnegie Mellon University Aug. 2026 – Present Graduate Student Researcher, advised by Prof. Changliu Liu Pittsburgh, PA Whole-Body and Safe Locomotion Research (ongoing): Investigating whole-body control, safety with control barrier functions (CBFs), and world-model-based locomotion. Dexmate Bimanual Manipulation Demo, IROS 2026: Contributed to a bimanual block-stacking demonstra- tion combining agent planning with policy execution. Implemented the agent component, prepared the on-site environment, adapted the system to different grippers, and performed system calibration. Computational Autonomy and Robotics Lab, UM Jun. 2025 – Apr. 2026 Research Assistant, Directed by Prof. Maani Ghaffari Ann Arbor, MI Online Learning Lie-MPC Field Perception: Built DINO/SAM2 segmentation and fine-tuned Grounded-DINO/YOLO for Robust ASV Control: Developed~~

a Neural-Fly-inspired learning-augmented Lie-MPC framework for an autonomous surface vehicle, estimating residual wind/current disturbances online and injecting learned dynamics corrections into the controller for robust trajectory tracking under hydrodynamic uncertainty. Implemented and evaluated the framework on a real ASV, achieved a 20% improvement [submerged vegetation detection](#) in trajectory-tracking performance over the baseline MPC controller during field experiments. Field Perception for Aquatic Vegetation Tracking: Built a water-scene perception pipeline for autonomous surface vehicles, combining DINO/SAM2-based segmentation with Grounded-DINO and YOLO fine-tuning to detect submerged vegetation and provide perception signals for navigation and autonomy experiments. ASV Hardware and Sensor Integration for Closed-Loop Field Autonomy: Designed sensor mounts and integrated GPS, lidar, IMU, and camera systems on an autonomous surface vehicle. Aligned and calibrated sensors to support field data collection and control experiments. Intelligent Robotics and Autonomy Lab, UM Oct. 2024 – Apr. 2026 Research Assistant, Directed by Prof. Vasileios Tzoumas Ann Arbor, MI Multi-Robot [Active-SLAM Simulation & ROS2 Integration: and ROS2](#): Rebuilt a Unity + AirSim multi-drone [simulation](#) testbed with [redesigned vehicle revised](#) dynamics and sensing modules, then [sensing](#); migrated SlideSLAM communication [logics](#) from ROS1 to ROS2 by refactoring topic interfaces and workflows for bandwidth-limited active-SLAM evaluation. Resource-Aware Multi-Robot Exploration Algorithms: Developed distributed exploration EXP3-style algorithms for bandwidth-limited robot teams, formulating communication-aware action selection with submodular objectives and bandit feedback in simulation to balance map coverage and communication cost. Surgical Computer Vision Detection, SJTU Apr. 2024 – Dec. 2024 Research Assistant, Directed by Prof. Yutong Ban Shanghai, China Foundation Segmentation Model Evaluation for Surgical Video: Evaluated Segment Anything Model variants on surgical video frames, analyzing domain shift and temporal mask stability around instruments, tissues, and tool-tissue interaction regions.

Publications and Manuscripts

Xiaohan Wang*, Ziyu Xu*, Xuyi Yang. “Re-encoding What You Keep After Every Compression: Retained Visual KV Is Computed in a Context That No Longer Exists”, Submitted to ICLR 2027. * Equal contribution. Xuyi Yang, Xiaohan Wang, Wenhao Zhang, and Ziyu Xu. “StreamJIT: Continuous Observation and Just-in-Time Deliberation for Proactive Streaming Video [Assistants](#)”, [Assistants](#)”, Submitted to ICLR 2027. [Yinan Dong, Ziyu Xu, Tsimafei Lazouski, Sangli Teng, and Maani Ghaffari](#). “Online Learning-Enhanced Lie Algebraic MPC for Robust Autonomous Surface Vehicle Control”, [IEEE Robotics and Automation Letters \(RA-L\)](#), [revised manuscript under review, 2025](#). C. Yuan, J. Jiang, K. Yang, Z. Xu, . . . , and Y. Ban. “Systematic Evaluation and Guidelines for Segment Anything Model in Surgical Video Analysis”, [npj Digital Surgery](#), accepted, 2024. [Yinan Dong, Ziyu Xu, Tsimafei Lazouski, Sangli Teng, and Maani Ghaffari](#). “Online Learning-Enhanced Lie Algebraic MPC for Robust Autonomous Surface Vehicle Control”, [IEEE Robotics and Automation Letters \(RA-L\)](#), [revised manuscript under review, 2025](#). [Selected](#)

Awards

Gold [Medal in 2023 Medal](#), University Physics Competition, by American Physical Society & American Astronomical Society (top 1%) Third place in 2023 RoboMaster Campus Competition (top 10%) [\(2023\)](#); Dean’s Honor List, University of Michigan College of Engineering (2024, 2025) M Prize in 2024 Mathematical Contest in Modeling (top 13%) [2025](#)).

Skills

Robot Learning & Control: VLA post-training, humanoid foundation models, whole-body policy, RL, real-robot policy deployment, teleoperation/data collection, MPC, sim-to-real Perception & Multimodal ML: PyTorch, OpenCV, [visual KV compression](#), [streaming/multi-turn inference evaluation](#), [LLM fine-tuning](#), SAM2, Grounded-DINO, YOLO, [DINO](#) features, visual token compression, spatial grounding, LLMs fine-tuning [OpenCV](#) Systems & Tools: ROS/ROS2, [Systems](#): Python/C++, [Unity + AirSim](#), [Docker](#), Linux, Git, [sensor integration](#), CAD, FEA, Manufacturing, embedded system [Docker](#), ROS/ROS2, [real-robot policy deployment](#)